

# Implementation of simple microprocessor using verilog

Asked 2 years, 3 months ago   Active 2 years, 3 months ago   Viewed 2k times



0



I am trying to make a simple microprocessor in verilog as a way to understand verilog and assembly at the same time.

I am not sure if I am implementing what I think of microprocessors well enough or if I am completely wrong. Should I simplify the idea of a microprocessor or I should program it as I am making one using real chips. for example, should I define a variable called address and make a big case statement that takes the assembly commands and does stuff with the memory and addresses. So far I have done something similar to this.

```
case (CMD_op)
  //NOP
  4'b0000: nxt_addr = addr + 4'b0001 ;
  //ADD
  4'b0001: begin
    op3_r = op1_r + op2_r;
    nxt_addr = addr + 4'b0001;
  end
```

CMD\_op is a 4-bit input that refers to a set of predefined 16 commands inside the case statement I have added above, that's just the first two cases, I have made a case for each command and how it tamper with the address. I have a 16-bit x 16-bit array that should hold the main program. The first 4-bits of each line referring to the assembly command and the next 12-bits referring to the arguments of the command.

for example here is the unconditional jump command JMP

```
//JMP
4'b0101: nxt_addr = op1_r ;
```

the 4'b0101 is a case inside the commands' case statement.

The reason I am asking this question because I feel that I am emulating a microprocessor rather making one, I feel like I am just emulating what a specific assembly command would do to the internal registers inside the microprocessor. I don't have a bus, but what would a bus do if I can skip its uses using Verilog.

I feel something is missing, Thanks.

[assembly](#) [verilog](#) [microprocessors](#)

asked Jul 30 '18 at 10:46

[Adoobli](#)



131 1 6

Have you studied the structure of a simple 5 stage microprocessor? e.g. DLX? – Qazi Jul 30 '18 at 10:52

@Qazi no, in fact I haven't taken the microprocessors course in EE yet, and I have started playing with verilog around 1 month ago – Adoobli Jul 30 '18 at 11:03

- 1 Without knowing exactly every stage of a simple pipeline and what that stage does, it is highly unlikely that your processor will look like a real one. Playing with code is fun but in Digital Hardware it is very important to know the theory behind the design precisely. I think if you study a simple DLX processor, it wouldn't take you more than a week. The tasks are divided into fetch, decode, execute, mem, and writeback stage. You need to study these and write one module / stage + memories to make a nice hobby processor. I did the same project as a part of my degree. I also made an assembler. – Qazi Jul 30 '18 at 11:09

- 1 Here is my design: [researchgate.net/publication/...](https://researchgate.net/publication/...) – Qazi Jul 30 '18 at 11:10

@Qazi thanks for the suggestion, I was thinking about making an assembler later on, but I didn't want to proceed until I knew what I was doing is right or wrong, I don't know what a DLX is, I just read a simple book describing microprocessors and I just summarized it without expanding the idea and implementing every part of it because I thought summarizing it will work. I haven't finished what I am working on, but in theory, I believe what I am doing will work, but it will not honour the idea of a microprocessor, I have a colourful history of awesome stuff that didn't end well. – Adoobli Jul 30 '18 at 11:26

## 1 Answer

Active	Oldest	Votes
--------	--------	-------



6



As detailed in the comments, there seems to primarily be confusion over how to handle memory/bus as well as some general questions over how to implement things across modules.

While SO isn't well designed to answer these broad questions of design/implementation of generic single-cycle processor, I will go through the steps of a VERY basic one here as a brief tutorial to clarify some point the author has.

### Step 1: ISA

First, the Instruction Set Architecture must be known and specified what each instruction does. Things in the ISA are the instructions themselves, the number of registers in the system, how interrupts and exceptions are handled among other things. Usually, engineers will use a preexisting instruction set (x86, ARM, MIPS, Sparc, PowerPC, 68000 etc) rather than designing a new one from scratch, but for learning purposes, I'll design our own. In the case I will show here, there will only be 4 basic instructions: LD (Load data from memory into register), ST (Store data from memory into register), ADD (Add registers together) and BRZ (Branch if last operation was equal to zero). There will be 4 General Purpose Registers and a Program Counter. The processor will do everything in 16 bits (so a 16-bit word). Each instruction will be broken down like so:

[15 OPCODE 14] | [13 SPECIFIC 0] -- Opcode is always in the top two bits, the rest of the instruction depends on the type it is

ADD: add rd, rs1, rs2 -- rd = rs1 + rs2; z = (rd == 0)  
 [15 2'b00 14] | [13 rd 12] | [11 rs1 10] | [9 rs2 8] | [7 RESERVED 0]

LD: ld rd, rs -- rd = MEM[ra]  
 [15 2'b01 14] | [13 rd 12] | [11 ra 10] | [9 RESERVED 1] | [0 1'b1 0]

```

ld rd, $addr -- rd = MEM[$addr]
[15 2'b01 14] | [13 rd 12] | [11 $addr 1] | [0 1'b0 0]

ST: st rs, ra -- MEM[ra] = rs
[15 2'b10 14] | [13 RESERVED 12] | [11 ra 10] | [9 rs 8] | [7 RESERVED 1] | [0 1'b1 0]

    st rs, $addr -- MEM[$addr] = rs
[15 2'b10 14] | [13 $addr[10:7] 10] | [9 rs 8] | [7 $addr[6:0] 1] | [0 1'b0 0]

BRZ: brz ra -- if (z): pc = ra
[15 2'b11 14] | [13 RESERVED 12] | [11 ra 10] | [9 RESERVED 1] | [0 1'b1 0]

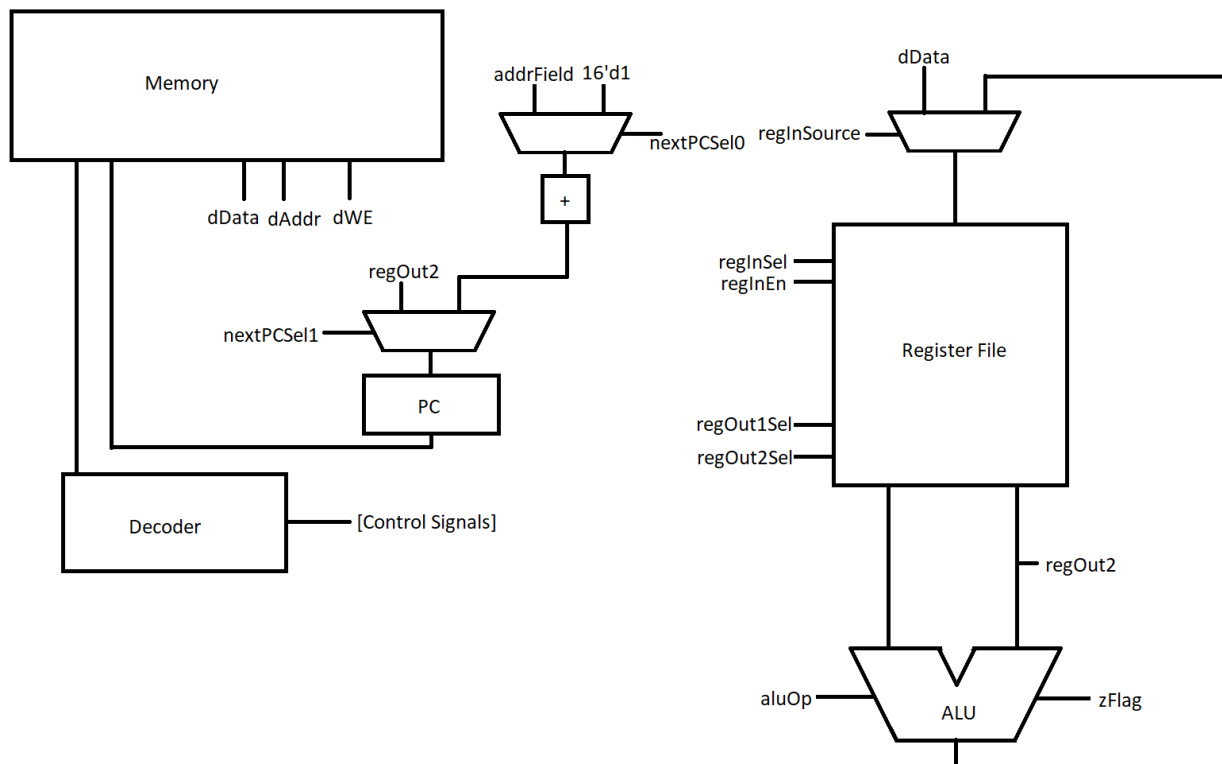
    brz $addr -- if (z): pc = pc + $addr
[15 2'b11 14] | [13 RESERVED 12] | [11 $addr 1] | [0 1'b0 0]

```

Note that there's different flavors of many instructions as a result of different ways to address memory (LD / ST both allow for register addressing and absolute addressing); this is a common feature in most ISAs, a single opcode might have additional bits that specify more details on the arguments.

## Step 2: Design

Now that we have an ISA, we need to implement it. To do so, we will need to sketch out the basic building blocks of the system. From the ISA, we know this system requires a 4x16-bit register file (r0 - r3) and register for pc (program counter), a simple ALU (Arithmetic Logic Unit, in our case it can only add) with Zero status register (z flag) and a bunch of combinational logic to tie to all together (for decoding the instructions, determining the next value of pc, etc). Typically, actually drawing it all out is the best approach, making it as detailed as needed to specify the design. Here is it done in some detail for our simple processor:



Notice the design is a bunch of the building blocks discussed before. Included also are all the of data lines, control signals and status signals in the processor. Thinking through everything you need before going to code is a good idea, so you can more easily modularize your design (each block can be a module) and see any major challenges beforehand. I want to note that I did notice a few mistakes/oversights on this diagram while going through implementation (mostly in missing details), but its important to note that the diagram is a template for whats being made at this point.

### Step 3: Implement

Now that the overall design is complete, we need to implement it. Thanks to having drawn it out in detail before hand, that just comes down to building up the design one module at a time. To start, lets implement the ALU as its pretty simple:

```

        input rst,
        input [15:0] in1,
        input [15:0] in2,
        input op, // Adding more functions to the system means adding bits to this
        output reg [15:0] out,
        output reg zFlag);

reg zFlagNext;

// Z flag register
always @(posedge clk, posedge rst) begin
    if (rst) begin
        zFlag <= 1'b0;
    end
    else begin
        zFlag <= zFlagNext;
    end
end

// ALU Logic
always @(*) begin
    // Defaults -- I do this to: 1) make sure there are no latches, 2) list all
    variables set by this block
    out = 16'd0;
    zFlagNext = zFlag; // Note, according to our ISA, the z flag only changes when an
    ADD is performed, otherwise it should retain its value

    case (op)
        // Note aluOp == 0 is not mapped to anything, it could be mapped to more
        operations later, but for now theres no logic needed behind it
        // ADD
        1: begin
            out = in1 + in2;
            zFlagNext = (out == 16'd0);
        end
    endcase
end

endmodule

```

To address your concern about behavioral Verilog; yes, you are writing code that is higher level and might seem like emulation. However, when doing Verilog, you are really implementing a hardware design. So, while you might write a line like `out = in1 + in2`, recognize you are actually instantiating an adder in the design.

Now, lets implement the register file:

```

module registerFile(input clk,
                   input rst,
                   input [15:0] in,      // Data for write back register
                   input [1:0] inSel,     // Register number to write back to
                   input inEn,           // Dont actually write back unless asserted
                   input [1:0] outSel1,   // Register number for out1
                   input [1:0] outSel2,   // Register number for out2
                   output [15:0] out1,
                   output [15:0] out2);

    reg [15:0] regs[3:0];

    // Actual register file storage
    always @(posedge clk, posedge rst) begin
        if (rst) begin
            regs[3] <= 16'd0;
            regs[2] <= 16'd0;
            regs[1] <= 16'd0;
            regs[0] <= 16'd0;
        end
        else begin
            if (inEn) begin // Only write back when inEn is asserted, not all instructions
                write to the register file!
                regs[inSel] <= in;
            end
        end
    end

    // Output registers
    assign out1 = regs[outSel1];
    assign out2 = regs[outSel2];

endmodule

```

See how we can treat every big block in our design diagram as a separate module to help modularize the code (literally!), so it separated functional blocks into different parts of the system. Notice also that I try to minimize the amount of logic inside the `always @(posedge clk)` blocks. I do this as its generally a good idea to understand whats a register and whats combinational logic, so separating them in the code helps you understand your design and the hardware behind it, as well as avoiding latches and other problems synthesis tools might have with your design when you get to that stage. Otherwise, the register file shouldnt be too surprising, just a "port" for writing back a register after an instruction is run (like `LD` or `ADD` ) and two "ports" for pulling out register "arguments."

Next is memory:

```

module memory(input clk,
              input [15:0] iAddr, // These next two signals form the instruction port
              output [15:0] iDataOut,
              input [15:0] dAddr, // These next four signals form the data port
              input dWE,
              input [15:0] dDataIn,
              output [15:0] dDataOut);

    reg [15:0] memArray [1023:0]; // Notice that Im not filling in all of memory

```

with the memory array, ie, addresses can only from \$0000 to \$03ff

```

    initial begin
        // Load in the program/initial memory state into the memory module
        $readmemh("program.hex", memArray);
    end

    always @(posedge clk) begin
        if (dWE) begin // When the WE line is asserted, write into memory at the
given address
            memArray[dAddr[9:0]] <= dDataIn; // Limit the range of the addresses
        end
    end

    assign dDataOut = memArray[dAddr[9:0]];
    assign iDataOut = memArray[iAddr[9:0]];

endmodule

```

A few things to note here. First, I kinda cheat a bit and allow for combinational memory reads (the last two `assign` statements), ie theres no register on the address and data lines of the memory array as there would be in most actual hardware (this design is probably going to be expensive on an FPGA). Its important to understand what kind of hardware your design will be synthesized into to avoid long combinational chains or impractical memories. Note also that the memory doesnt fill the entire  $2^{16}$  possible address space. Its not common in computer systems to have as much physical memory as the address space allows for. This opens up those memory addresses for peripherals and other memory mapped IO. This is generally what you might call the system's bus, the interconnect between memory, the CPU and any other peripherals. The CPU accessed the bus via its instruction read port and its data read/write port. In this system, the memory used for storing instructions and data is the same, so called von Neumann architecture. Had I separated the instruction memory from the data memory (ie, two separate memory modules), it would be a Harvard architecture.

Onward to the final submodule, the instruction decoder:

```

    dAddrSel = 1'b0; // Choose to use addr as dAddr
    dWE = 1'b1; // Write to memory
    addr = {6'b0, instruction[13:10], instruction[7:1]}; // Zero fill addr to get
full address
    end

    // Register
    1'b1: begin
        dAddrSel = 1'b1; // Choose to use value from register file as dAddr
        dWE = 1'b1; // Write to memory
    end
    endcase
end

// BRZ
2'b11: begin
    // Instruction does nothing if zFlag isnt set
    if (zFlag) begin
        // BRZ has 2 versions, register addressing and relative addressing, case on
that here
        case (instruction[0])
            // Relative

```

```

1'b0: begin
    nextPCSel = 2'b01; // Select to add the addr field to PC
    addr = {{6{instruction[11]}}, instruction[11:1]}; // sign extend the addr
field of the instruction
end

// Register
1'b1: begin
    nextPCSel = 2'b1x; // Select to use register value
end
endcase
end
end
endcase
end

endmodule

```

In the design I provided above, each module had a number of control signals (like the memory `dWE` to enable memory writes on the data port; `regSelIn` to select the register in the register file to write to; `aluOp` to determine what operation the ALU should perform) and a number of status signals (in our design, that's just `zFlag`). The decoder's job is to take the instruction apart and assert the needed control signals based on what the instruction is trying to do, sometimes with help from the status signals (like how `BRZ` needs `zFlag`). Sometimes, the instruction itself encodes these signals directly (like how `regInSel`, `regOutSel1` and `regOutSel2` can be pulled out of the instruction word itself) but other times these control signals do not map directly (like `regInEn` doesn't really map to any single bit in the instruction word).

In your design, it seems like you were doing a lot of the actual work of the instructions inside your decoder itself, and that's fine sometimes, but it usually will result in a bunch of extra hardware (ie, similar instructions will not share hardware, like an increment instruction and add instruction will not share an adder typically in your coding style, but they should in a real design). Separating the system into a control path and data path, where the control path asserts control signals to instruct the data path how to handle data, while the data path does the actual work and returns status signals to indicate anything important.

The final step is to bring it all together and add in the parts of the hardware that didn't neatly fit into a nice box (like the program counter, don't forget that!):

```

    .addr(sel(addrSel)),
    .addr(addr));

// PC Logic
always @(*) begin
    nextPC = 16'd0;

    case (nextPCSel)
    // From register file
    2'b1x: begin
        nextPC = regOut1;
    end

    // From instruction relative
    2'b01: begin
        nextPC = PC + addr;
    end
end

```

```

// Regular operation, increment
default: begin
    nextPC = PC + 16'd1;
end
endcase
end

// PC Register
always @(posedge clk, posedge rst) begin
    if (rst) begin
        PC <= 16'd0;
    end
    else begin
        PC <= nextPC;
    end
end

// Extra logic
assign regIn = (regInSource) ? dDataOut : aluOut;
assign dAddr = (dAddrSel) ? regOut1 : addr;

endmodule

```

See that my processor is now just a bunch of module instantiations and a bit of extra registers and muxes to link it all together. These do add a few extra control signals to our design though, so make sure you think it out a bit as part of the overall system design. Its not a big detail however to go back and add these new signals to the decoder, but youll probably have already realized you needed them at this point! One other thing to note is its not typical to include memory in the processor itself. As mentioned before, memory is separate from the CPU and these two are typically connected together outside the processor itself (so, should be done outside the processor module); but this is a quick and simple introduction so Im putting it all here to avoid having to have another module that includes the processor and memory and connects them together.

Hopefully this practical example shows you both all the steps and all the major components and how to implement them. Note that I didnt full validate this design, so its possible I made a few mistakes in the code (I did run a few tests though so it should be ok :) ). Again, this kind of thing isnt the best for SO, you should ask specific questions as broad topic questions are typically closed quickly. Note also that this is a brief and SUPER simple introduction, you can find more online and theres ALOT more depth to computer architecture than this; pipelining, interrupts/exceptions, caching all come to mind as next topics. And this architecture doesnt even have any kind of stalling for memory, not multiword fetching for instructions and alot more common things you find in even the smallest processors.

edited Aug 1 '18 at 4:54

answered Jul 31 '18 at 19:58



Unn

3,723 10 25

---

Hi, reading that was a blast and I will read it many times after this to make sure I haven't missed anything. I want to ask why did you make `in1` and `in2` in the `ALU` module 16-bit inputs while they can only be 4 bits when used as you mentioned in the `ADD` instruction example at the beginning. I have made a pastebin link to my Verilog code, it's far from over, I just wrote the function of each instruction inside the decoder, so it's far from done. [my code](#), I have even included a 4-bit error register which shows an error with instruction misuse – [Adoobli](#) Jul 31 '18 at 21:51




Thank you very much for the detailed answer. There are few stuff to note of what I think of microprocessors. There is an awesome youtuber called [Ben Eater](#) who did a full series on how he made a processor and I watched it and got a general idea from it. As for the memory, I didn't think of it until you mentioned it, what my `load` command does is it just stores whatever argument it gets in the last two addresses in my program, my main array is a 16x16 array and I made it possible to store a value in the last two addresses, I didn't -- -- [AdoobII](#) Jul 31 '18 at 21:56

---

@AdoobII The ALU is the module that actually does the operation and thus needs the full operands (ie, all 16-bits). One way to think of the design is as a road for data words (words in this case being 16-bits), ie the data path. The control path is like stop lights that can be green or red, selecting which path the data can take and from where it comes (ie, selects what register goes out `out1` from the register file, and where that data goes). That YouTube channel is interesting, but I think that guy is a level below what you need to worry about if you are writing behavioral Verilog. -- [Unn](#) Jul 31 '18 at 22:01

---

-- think of external memory, I just thought of translating a C code into assembly in my mind. The `ADD` and the other operations that requires two arguments only take immediate values, because I forgot about the ability of making variables or storing something at all, note that I mixed store with load, so my load is technically a store, but it stores an immediate value or a result from the instruction before it into an address inside the program itself. As I am writing this comment I've realised that I've missed many stuff but as I said it's far from complete, Do you recommend continuing -- -- [AdoobII](#) Jul 31 '18 at 22:01 

---

-- the abomination I am making? -- [AdoobII](#) Jul 31 '18 at 22:02

---